


LLMs in technical writing

The fundamentals of technical writing | VUT 2025

The Red Hat Customer Content Services team

About the authors:



Alexandra Nikandrova
Technical Writer for Red Hat
Former Devops



Jana Vrbkova
Technical Writer for Red Hat
Favorite AIs: Google AI on Pixel
DeepL




Srikanth R
Technical Writer for Red Hat

What we'll discuss today

- What is LLM and SLM?
- Types of LLM
- Use cases of LLM in technical writing
- LLM vs automation
- Pros and cons of LLM
- Prompt engineering






How do large language models work?

- LLMs have trillions of parameters.
- How LLMs work:
 - **Data**
 - **Architecture**
 - **Training**
- They can be **fine-tuned** for specific tasks.
- They iteratively **predict the next token** based on context.

Types of LLMs



- Architecture (encoder, decoder, mixture of experts)
- Training methodology (general purpose LLM, fine-tuned)
- Modality (text-only, multimodal)
- Licencing (proprietary, open-source)

Types of LLMs

Encoder, decoder, mixture of experts

Encoder-only focus on understanding input data

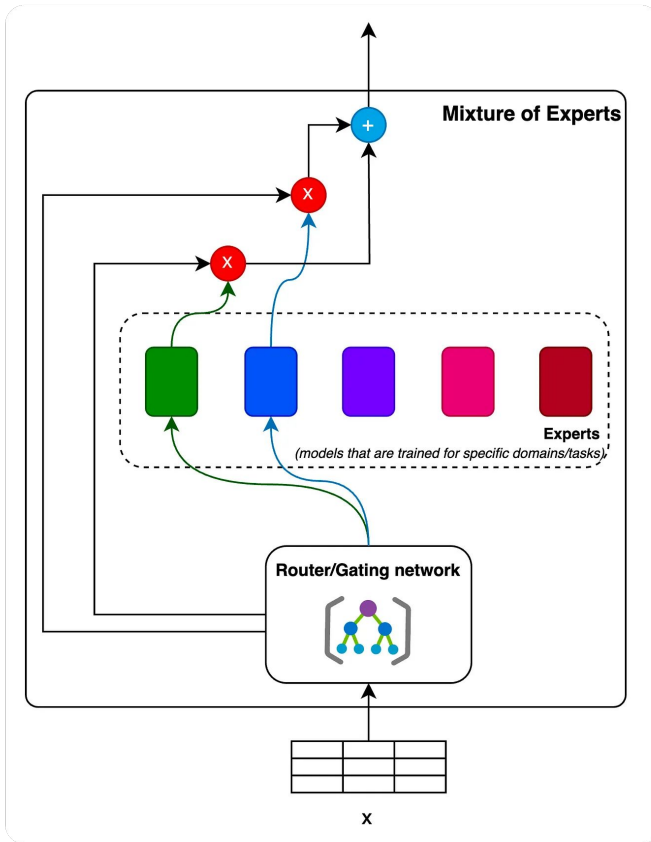
Examples: BERT (Bidirectional Encoder Representations from Transformers)

Decoder-only focus on generating context-aware text

Examples: GPT, LLaMA, Gemma, Claude

Mixture of Experts (MoE) focus on scalability and efficiency of all the models involved

Examples: Qwen3 (30B A3B)



Types of LLMs

General purpose vs Fine-tuned



General purpose LLMs are trained on a wide range of data—can handle most tasks but not optimized for one domain.

Examples: ChatGPT, Claude, Gemini

Fine tuned LLMs are Specialized on specific domains like legal, medical, or technical documentation.

Examples: Writer.com (custom LLM), CodeLlama (for code), BioGPT (for biomedical content)


Types of LLMs



Retrieval-Augmented Generation (RAG)

General purpose model fine-tuned to your needs.

Types of LLMs



Text-only vs Multi-modal

Text-only LLMs are primarily designed to understand, analyze, and generate text-based content.

Examples: GPT-3, LLaMA, BLOOM.

Use Cases: Text generation, translation, summarization, question answering, chatbots.

Multi-modal LLMs are capable of processing and generating content across multiple data types, such as text, images, audio, and video.

Examples: Flamingo (DeepMind), Kosmos-1 (Microsoft), LLaVA (open-source), PaLM-E (Google).

Use Cases: Computer Vision, Audio-Visual AI, Robotics, Content Creation, and so on.

Types of LLMs


Open-source vs Proprietary


Open-source LLMs are publicly available, allowing for free use, modification, and distribution, enabling extensive customization.

Examples: Llama 3, Falcon, WizardLM, and so on.

Proprietary LLMs have limited access and customization options, often offered with strong performance and robust support.

Examples: OpenAI's GPT-4, Google's Gemini, Anthropic's Claude, and so on.






How do you use AI tools?



LLM use cases in TW

- Brainstorming: “Give me 5 headings for this paragraph”
- Research, explanations, summarisation: “How does SELinux work?”
- Missing prerequisites: “What do I need to complete this procedure”
- Description of code: “What does this yaml file do?”
- Code snippets or regex: “Create a script that...”
- Formatting: “Convert this csv into a markdown table”
- Ideas for visual elements (runwayml.com)



Automation vs AI




When to use automation

- Predictable, repetitive rules (processing files, system monitoring)
- Low complexity
- Speed
- Limited data (no data required)
- Low cost (generally cheaper to implement)



When to use AI

- Pattern recognition with no clear rules
- Processing natural language
- Modeling



Pros and cons

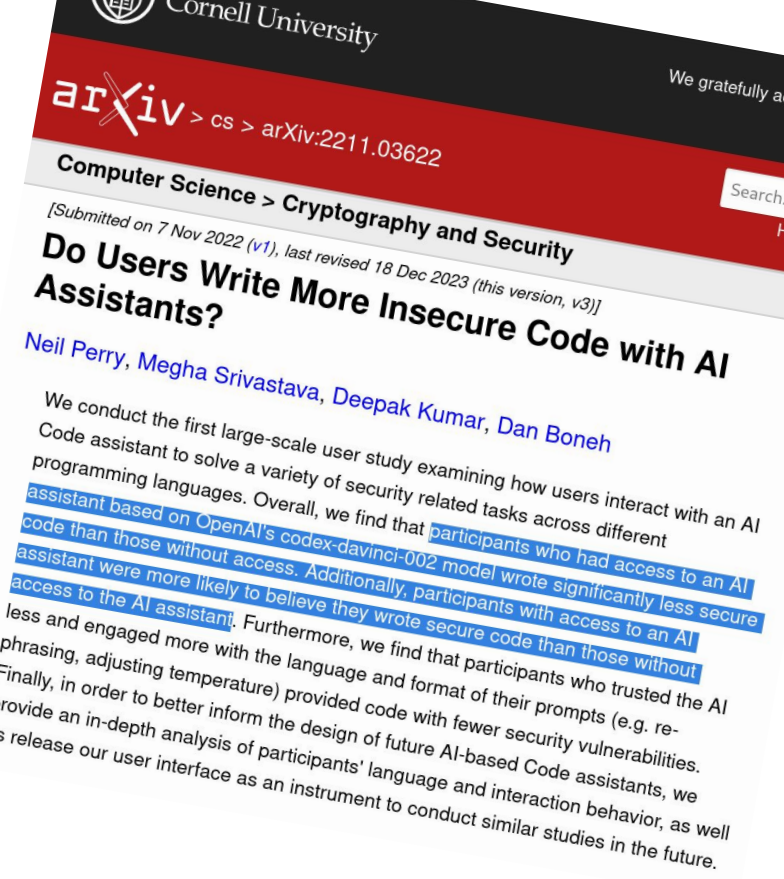
Pros of LLMs

- Enhanced efficiency
- Automation
- Reduced human error
- Reduced physical risks
- Improved customer experience
- Democratization of skills



Cons of LLMs

- No concept of truth
- Hallucinations
- Legal limitations
- Security and privacy concerns
- Environmental impact
- De-skilling



Sources: unsplash.com/@jannerboy62

<https://arxiv.org/abs/2211.03622>

<https://dl.acm.org/doi/10.1145/3491101.3519665>

<https://www.technologyreview.com/2025/05/20/1116327/ai-energy-usage-climate-footprint-big-tech/>

<https://www.nature.com/articles/s41586-024-07566-y#Sec2>


<https://www.sciencedirect.com/science/article/abs/pii/S0007681324000272>



Prompt Engineering

Which text is written by a human TW?

A	B	C
<p>Red Hat Enterprise Linux 9: GIMP 2.99.8 Technology Preview Available</p> <p>This release introduces GNU Image Manipulation Program (GIMP) version 2.99.8 as a Technology Preview in Red Hat Enterprise Linux 9. GIMP 2.99.8 delivers a comprehensive set of image editing and manipulation capabilities, enabling users to:</p> <ul style="list-style-type: none">• Perform basic image editing tasks: Crop, resize, adjust color and brightness, apply filters, and more.• Create and manipulate artwork: Draw vector graphics, paint with various brushes, and work with layers for complex compositions.• Automate repetitive tasks: Utilize scripting capabilities to automate workflows and save time. <p>Important Note: GIMP 2.99.8 is provided as a Technology Preview. This means it is still under development and may contain bugs or stability issues. It is not intended for production use.</p> <p>This Technology Preview allows users to explore the capabilities of GIMP 2.99.8 before its official introduction in Red Hat Enterprise Linux 9. Feedback from users will be crucial in shaping the final offering.</p> <p>Installation:</p> <p>GIMP 2.99.8 is available through the Technology Preview repository. You can install it using the following command:</p> <pre>sudo dnf install gimp-techpreview</pre> <p>Disclaimer:</p> <p>Red Hat does not provide support for Technology Preview features. Use them at your own risk and discretion.</p>	<p>Feature: GNU Image Manipulation Program (GIMP) 2.99.8</p> <p>Description: The GIMP package version 2.99.8 is now available as a Technology Preview in Red Hat Enterprise Linux 9. This pre-release version provides a glimpse into the upcoming GIMP 3 release. Users can explore and test this version using the RPM package for installation.</p> <p>Key Points:</p> <ul style="list-style-type: none">• Technology Preview Status:<ul style="list-style-type: none">◦ GIMP 2.99.8 is presented as a Technology Preview, showcasing features and enhancements that will be fully integrated into the official GIMP 3 release in future versions of RHEL 9.• Pre-Release Version:<ul style="list-style-type: none">◦ This package serves as a pre-release version, allowing users to experience and provide feedback on the latest developments in GIMP. <p>Recommended Action:</p> <ul style="list-style-type: none">• Users interested in previewing the upcoming features in GIMP 3 can install the GIMP 2.99.8 package.• As a Technology Preview, this version is not intended for production use but offers users an opportunity to familiarize themselves with the upcoming GIMP release. <p>Note: GIMP 2.99.8 is a forward-looking Technology Preview, giving users early access to the evolving capabilities of GIMP 3. Feedback from the community during this preview phase is valuable in refining and enhancing the user experience before the official release.</p>	<p>GIMP available as a Technology Preview in RHEL 9</p> <p>GNU Image Manipulation Program (GIMP) 2.99.8 is now available in RHEL 9 as a Technology Preview. The <code>gimp</code> package version 2.99.8 is a pre-release version with a set of improvements, but a limited set of features and no guarantee for stability. As soon as the official GIMP 3 is released, it will be introduced into RHEL 9 as an update of this pre-release version.</p> <p>In RHEL 9, you can install <code>gimp</code> easily as an RPM package.</p>



How to write prompts for LLM?

- Start simple
- Instruct
- Specify and direct
- To do



Takeaways

- LLMs are part of a foundational model
- Requires a lot of data and computing power to train
- Automation can be used cheaper and faster in some cases
- **AI cannot think**, only predict the next token. You need to have the following skills: user focus, minimalism, stylistics, applying a style guide



Exercise

gemini.google.com



- Access documentation
- Download a PDF file



Learn prompt engineering

[Prompt engineering guide](#)

EXERCISE: Now you try it!

Cause: Too restricted checks in openssl

Consequence: random tls connection failures

Workaround (if any): just reconnect

Result: handshake succeeds

gemini.google.com

bing.com/chat

chat.openai.com

Exercise: Tips for prompting

"Document how to manage and configure passkey authentication devices in the Identity Management (IdM) environment."

- Provide contextual details.
- Be clear and specific in your prompts.
- Focus on one aspect at the time (content format, structure, language).
- Work one section at a time (heading, abstract, procedural steps, verification...).
- Provide examples of the style or tone you're aiming for.
- Refine the prompts using different wording to get better results.
- Feed ChatGPT with content or reference material.

Exercise: Tips for prompting

Bonus: Custom instructions (system prompt)

Go to **Your profile > Customize ChatGPT** to customize the responses.

For example, create your own style guide:

1. **Be Clear:** Use simple and direct language. Avoid unnecessary jargon unless it's needed.
2. **Stay Formal:** Keep a formal and neutral tone. Don't add personal opinions or emotions.
3. **Talk to the Reader:** Address the reader directly with "you".
4. **Use Present Tense:** Describe actions or states in the present tense, not future tense.
5. **Use AsciiDoc:** Format text using AsciiDoc markup, not Markdown.
6. **Give Actions:** Provide clear instructions or guidance on what the reader should do.



Technical writing is not just writing

- Research of new topics
- Planning, content strategy, user/content journey
- Verification and testing
- User experience

Thank you